



Center for Visual Information Technology
IIIT Hyderabad

Reinforcement Learning - Introduction

Mohit Rathore
mrmohitrathoremr@gmail.com

Sep 7, 2018



Trivia

- Motivation

- News

Deep Q Learning

- Understanding the terms.

- The approach.

- Flaws with DQN.

Practical examples

- Gym and fromscratchtoml

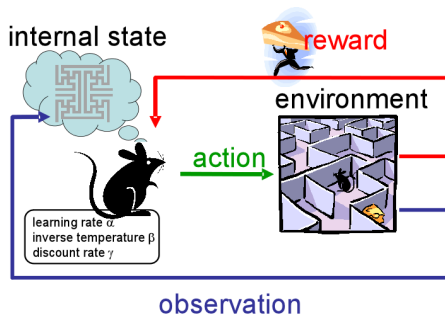
Why should I care?



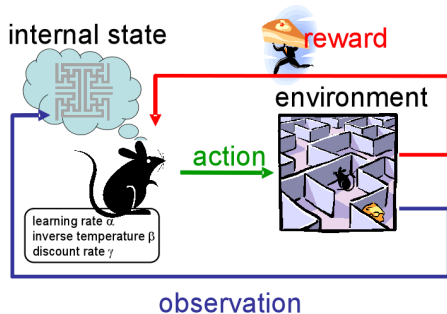
-

- ▶ RL Beats Humans at Dota 2. [▶ Link](#)
- ▶ AlphaGo beats 18-time world champion Lee Sedol. [▶ Link](#)
- ▶ Google has developed an approach to hyperparameter tuning using reinforcement learning that they call AutoML. [▶ Link](#)

- **State** - It can be seen as a variable which directly affect the action taking mechanism. In an environment such as a maze, a state can be seen as the x, y coordinates of the agent.




- ▶ **Action** - Performing a particular action in a particular state takes you to the next state.
- ▶ The mouse can move in either of the four directions to reach the next block.




- ▶ **Agent** - It is an abstraction of the entity which acts and learns from previous experiences in hope to maximize reward.
- ▶ In this case our agent is - mouse.



- ▶ **Reward** - It is the utility that our agent is trying to maximize.
- ▶ For our mouse reward is - cheese.

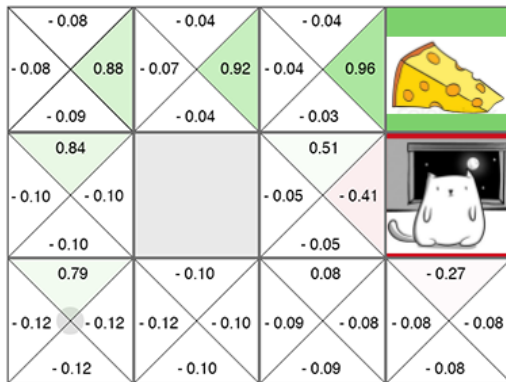


- **Environment** - Environment can be considered as an entity which takes in space-action and returns the next state and reward for that action.



Reinforcement Learning

Deep Q learning





What can be the best case scenario?

What can be the best case scenario?



Deep Q-Network Training

- Bellman Equation:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

- Loss function (squared error):

$$L = \mathbb{E}[\underbrace{(r + \gamma \max_{a'} Q(s', a'))}_{\text{target}} - Q(s, a))^2]$$



Why discount the reward by gamma?



Why discount the reward by gamma?

- So that rewards do not diverge as the time steps increase.

Why discount the reward by gamma?

- ▶ So that rewards do not diverge as the time steps increase.
- ▶ We will be running in circles.



Exploration!

- ▶ We take the optimum action but sometimes we take a random action with exploration probability ϵ so that surroundings are investigated for possible rewards.
- ▶ We decay this ϵ over time exponentially.

$$\epsilon = \epsilon_{min} + (\epsilon_{max} - \epsilon_{min}) * \exp^{(-\lambda * t)}$$



- **Explore** - Take random actions and create a queue of (current state, action, next state, reward).

Reinforcement Learning

Deep Q learning



- ▶ **Explore** - Take random actions and create a queue of (current state, action, next state, reward).
- ▶ **Train** - Sample from this queue and use the data to create a training data for our model to learn from.



- ▶ **Explore** - Take random actions and create a queue of (current state, action, next state, reward).
- ▶ **Train** - Sample from this queue and use the data to create a training data for our model to learn from.
 - ▶ This should be done keeping in mind that we tend to converge the bellman equation.




- ▶ **Explore** - Take random actions and create a queue of (current state, action, next state, reward).
- ▶ **Train** - Sample from this queue and use the data to create a training data for our model to learn from.
 - ▶ This should be done keeping in mind that we tend to converge the bellman equation.
 - ▶ Taking a random action with a exploration probability (which decrease over time).



- ▶ **Explore** - Take random actions and create a queue of (current state, action, next state, reward).
- ▶ **Train** - Sample from this queue and use the data to create a training data for our model to learn from.
 - ▶ This should be done keeping in mind that we tend to converge the bellman equation.
 - ▶ Taking a random action with a exploration probability (which decrease over time).
 - ▶ The new variables (state, action, next state, reward) acquired are pushed into the queue.



- ▶ **Explore** - Take random actions and create a queue of (current state, action, next state, reward).
- ▶ **Train** - Sample from this queue and use the data to create a training data for our model to learn from.
 - ▶ This should be done keeping in mind that we tend to converge the bellman equation.
 - ▶ Taking a random action with a exploration probability (which decrease over time).
 - ▶ The new variables (state, action, next state, reward) acquired are pushed into the queue.
- ▶ **Play** - The model is tested by keeping the exploration probability as zero and stopping learning from experiences.



Deep Q-Network Training

- Bellman Equation:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

There is a flaw in this approach.

There is a flaw in this approach.

- Our Q values shift but the target value also shifts.



There is a flaw in this approach.

- ▶ Our Q values shift but the target value also shifts.
- ▶ Due to the max in the formula for setting targets, the network suffers from maximization bias, possibly leading to overestimation of the Q function's value and poor performance.


$$\underbrace{Q(s, a)}_{\text{TD target}} = r(s, a) + \gamma \underbrace{Q(s', \argmax_a Q(s', a))}_{\substack{\text{DQN Network choose} \\ \text{action for next state}}}$$

Target network calculates the Q value of taking that action at that state



Like most all machine learning algorithms all this comes at a price.

Like most all machine learning algorithms all this comes at a price.



It uses a lot of system resources!

Breakout by deepmind [▶ Link](#)

Examples

openAI gym and fromscratchtoml



openAI gym provides an environment for your agent to learn and play around. [▶ Link](#)

Code for machine learning algorithms from scratch - [▶ Link](#)



Thank You